

I built an AI agent to audit a Databricks catalog.

34 → 89

That's the catalog health score, before and after one run. Everything on the next nine slides is what actually happened getting there — including the parts I got wrong.

A column holds the value

47382

Is that **47 seconds** or **47 minutes**? Both are plausible. One is wrong by a factor of a thousand.

You can profile that column all day — nulls, cardinality, samples. None of it answers the question, because the answer isn't in the data. It's in the vendor's documentation. That's the one thing a feature that only reads your data structurally cannot do.

Every description is generated **and grounded in the source system's own docs.**

79/79

columns documented

\$0.0048

modelled cost, list price

118s

for the whole catalog

Free tier bills nothing, so that figure is modelled, not billed — and the repo says so instead of quoting a number it didn't pay.

My first test **was rigged**, and it took me a day to notice.

I wrote the documentation.
I generated the data.
I declared the result.

That test could not fail. A test that can't fail proves nothing — it only proves text I put into a prompt came back out.

So I ran it again against material **neither side authored**.

The official **TPC-H v3.0.1** spec from tpc.org, against the sample tables that ship with Databricks. 61 columns.

clean sweep → 6 real fixes

"56 of 61 changed" sounds better and means almost nothing — that's just rewording. About 6 changed **materially**. Six out of 61, on data I didn't touch, is a real result. The clean sweep was a mirror.

The correct answer was already in its prompt. It ignored it.

Three TPC-H columns hold single letters — the exact derivation rule was retrieved and sitting in context.

Model wrote: "O for open, F for final"

A confident, familiar guess about a famous benchmark — beating the text it was actually given. The fix wasn't better retrieval. It was one paragraph telling it to use the rule instead of its own instinct.

To find abandoned tables, the tool has to read them.

Reading is the one thing "abandoned" means the absence of. Three attempts before the fix held:

01 - failed

Subtract the scan's time window

Unity Catalog logs events asynchronously — the tool's own reads landed outside their own window.

02 - failed

Pad the window by 20 minutes

Fixed that, and swallowed a real analyst session that happened 4 minutes after a scan.

03 - held

Mark every statement, match by identity

Not by timing. The agent became invisible to itself — verified in the same run that profiled every table.

**Every description, in English
and Portuguese — one
model call.**


\$0.0048 → \$0.0062

+29% for twice the coverage, because the expensive part
of the prompt is identical for both languages.

`wrapUpCode` stays `wrapUpCode`

`"F"` never becomes `"final"`

Not a slideshow. A portal you can actually open.



Click any table. Every description links to the exact passage of documentation that produced it.

What I'd still bet against — and what's already working.

16

personal-data columns
caught

2

tables truly never
read

70

tests, zero secrets
needed

Full write-up, every commit, the failed attempts included —
link in the first comment.